

# **++Malloc: Asst.1 - Systems Programming**

By: Kevin Shah (kas665) & Manav Patel (mpp124)

## **Introduction:**

The purpose of this assignment was to implement `malloc()` and `free()` in our own way using our knowledge about pointers, arrays, loops, casting, and other concepts of C programming. One of the main obstacles was to figure out how to work with header files, Makefiles, and redirecting all calls from `malloc()` and `free()` to `mymalloc()` and `myfree()` respectively using macros. After learning the basic concepts, our next objective was to understand how we can compile our functions in `mymalloc.c` and output it as an object `mymalloc.o` file. Putting all of these files together, we managed to create a Makefile that not only compiles, but also executes and cleans the files all in one command - “make”.

## **Design & Implementation:**

One of the most important tasks of this project was to design an implementation that would allow us to traverse through our limited static array of bytes and determine if a particular block of bytes were allocated or free. A possible way to do this would be to store “metadata” for every block that is allocated or freed and then use a linked list to traverse from one block to another. This would require us that we have an integer or a short that tells us if a particular block is free or allocated and store a pointer to the next block. Thus, putting both of these elements into a struct of a linked list type it is possible. However, this is a very **inefficient** way to go about creating the function. Instead, we used one short as the metadata for every block and traversed through the array by finding the metadata and figuring out how many bytes away the next block is. The problem with this implementation is we only know the size of the block but we do not know if it is free or not. The way we fixed this issue was by assigning:

- **Negative Value = Free Space**
- **Positive Value = Malloced Space**

Example:

-13 = 13 Bytes of free space

13 = 13 Bytes of allocated space

0 = Error (Would never happen)

Now that we had a working and efficient implementation of the metadata, all we had left to do was how we determine during our first cold run of malloc or free if it has a correct design structure or not. We did this by creating a “magic-number”: A random hard-coded number that would allow us to determine if our implementation exists or not.

Our final implementation looked something like this:

Magic Number	10	Data of 10 Bytes	-1	1 Byte Free Space	2	Data of 2 Bytes	.....
--------------	----	------------------	----	-------------------	---	-----------------	-------

## Malloc and Free Implementation

*A short synopsis on how the methods actually function.*

Mymalloc was implemented in the following way. First step was to check if it was initialized or not. It did this by casting the first two bytes of myblock to a short and checking for a specific number which we refer to as the magic number. If it doesn't find it, it initializes the first two bytes to the magic number. It also creates a metadata block holding the size of the first malloc call, and creates a second metadata block containing the information for the rest of free space. If it has been initialized, it traverses the myblock, jumping from one metadata block to the other, until it finds a block both big enough and free. After doing so, it uses the createmetadata method to again split the free block into two, creating metadata that holds the size of the malloced block, and creates another metadata (the amount malloced spaces down) that holds the size of the remaining free block that was split.

Myfree was implemented in the following way. The first step was to by using the `check()` function to find if the pointer was valid, within bounds, pointing to a the beginning of the memory block, not already freed. If it was, it continued on, and returned the metadata size of the previous block, to be used later. if not, it returned a slew of values which indicated what type of error it was. Next, using the `merge()` function, it checked the adjacent memory blocks to see if they were free and combining them. It would check the right one first, and if it was free, merge it into the left one. Next it would check the left block. If it was free, it would merge the current block into the left block.

### **Space & Time Efficiency:**

Our metadata size consisted of only two bytes. We wanted to make our implementation for `mymalloc()` and `myfree()` as space and time efficient as possible with a first-fit algorithm. So in order to do so, we found the best way to do this was by using shorts instead of structs or pointers because pointers would typically be 4 - 8 bytes long depending on the system and structs would be at least the size of all the data types in it and additional padding. Therefore, by making one short to store the magic number for initialization and check purposes at the start of the array, and with metadata of just one short for every block, we minimize the need to store much data. By using arithmetic calculations to traverse through the array we would be time efficient as well because we would not be searching through every single byte for something. In this way, the metadata served as a pseudo pointer to the next block. This especially worked because of the way `mymalloc` worked. Every free/occupied block is managed by metadata. And because of the robustness of the algorithm, there would never be a case in which the pointer used for traversal would be pointing at anything but metadata. Thus, there was no need to include a magic number at every metadata block, for we would only *be* traversing to metadata blocks. Since there was only one element to the metadata, there was no need to include a structure for it. Thus with this ingenious method of traversal, and using signs to keep track of free/occupied blocks, we were able to store the metadata in a meager two bytes.

**Workload Data:**

We tested all test cases A through F and intensively tested for any breaks or faults in our code 100 times and recorded the time for each iteration of the workload and calculated the average time of all workloads. The workload on memgrind.c was as follows:

A: malloc() 1 byte and immediately free it - do this 150 times.

B: malloc() 1 byte, store the pointer in an array - do this 150 times. Once you've malloc()ed 50 byte chunks, then free() the 50 1 byte pointers one by one.

C: Randomly choose between a 1 byte malloc() or free()ing a 1 byte pointer > do this until you have allocated 50 times - Keep track of each operation so that you eventually malloc() 50 bytes, in total > if you have already allocated 50 times, disregard the random and just free() on each iteration - Keep track of each operation so that you eventually free() all pointers > don't allow a free() if you have no pointers to free().

D: Randomly choose between a randomly-sized malloc() or free()ing a pointer – do this many times (see below) - Keep track of each malloc so that all mallocs do not exceed your total memory capacity - Keep track of each operation so that you eventually malloc() 50 times - Keep track of each operation so that you eventually free() all pointers - Choose a random allocation size between 1 and 64 bytes.

E: Allocate an array to maximum capacity in mallocs of 1 byte and then free each byte randomly.

F: Completely allocates the array to maximum capacity in chunks of 100 and then free random blocks, then reallocate memory in size of 33 and do this 150 times. Finally, free the rest of the space.

**Findings:**

To actually test and find the average run time of each test case, we ran each test case 10,000 times. This was because, sometimes the program ran so fast, it didn't register an actual time. Therefore, 10,000 runs was more accurate in figuring out the actual average.

Average Time Taken by Test Case A: 0.00000200 seconds

Average Time Taken by Test Case B: 0.00002700 seconds

Average Time Taken by Test Case C: 0.00000700 seconds

Average Time Taken by Test Case D: 0.00000600 seconds

Average Time Taken by Test Case E: 0.00860700 seconds

Average Time Taken by Test Case F: 0.00004600 seconds

Its quite clear, that our malloc/free implementation is time-efficient as well. Each workload except E, takes a fraction of a millisecond indicating the efficient speed at which our implementations functioned.

The interesting thing in these findings is the significant difference between our Test Case E and others. The reason why workload E took so long compared to the other workloads is because the array of 4096 was allocated 1 byte at a time to its maximum capacity and then freed random 1 byte allocations. The real time consuming part of the workload was that each pointer was freed randomly. So the function had to wait until a random number was generated for each of the numbers from 0 to 1364, and due to repeats and randomness, it took a long time to free the entire block.